

Majeti Lahari

laharimajeti2002@gmail.com | +91-8121438529

github.com/Lahari3769 | linkedin.com/in/majetilahari | majetilahari-portfolio.vercel.app

Professional Summary

Software Engineer with 3+ years of experience building and deploying production-grade GenAI, RAG, and AI-driven systems on cloud infrastructure. Proven ability to automate large-scale workflows, improve system reliability, and deliver measurable business impact. Experienced in LLMs, Agentic AI, CI/CD, cloud-native DevOps, and leading high-performing engineering teams.

Skills

Languages: Python, TypeScript, PHP (Laravel), SQL, Shell Scripting, HTML, CSS.

GenAI & ML: Large Language Models (Watsonx.ai, Hugging Face, OpenAI, Google Gemini), RAG, Agentic AI, LangChain, NLP, Machine Learning, Deep Learning.

Frameworks & Tools: Flask, Streamlit, Angular, Docker, Git, Postman, Linux.

Cloud & DevOps: IBM Cloud, OpenShift, CI/CD Pipelines (Jenkins, Tekton), Secrets Manager, Security Scanning (Mend, Twistlock), Cloud Object Storage.

Databases & Vector Stores: IBM Db2, Milvus, ChromaDB.

Professional Skills: Agile Delivery, Technical Leadership, Mentorship, Cross-functional Collaboration, Executive Communication.

Work Experience

IBM, Bangalore

Software Developer (DevOps)

Apr 2025 -

Present

- Architected automated dashboards providing real-time feeder and process visibility, eliminating manual tracking and enabling one-click operational insights for application and leadership teams.
- Designed and deployed a RAG-based production chatbot on IBM Cloud using Flask, Docker, Jenkins CI/CD, and Milvus vector search integrated with Watsonx.ai LLMs, enabling instant access to 100% of project documentation and reducing manual data retrieval by 90%.
- Implemented secure, ITSS-compliant CI/CD pipelines with Mend and Twistlock security scanning, publishing audit-ready reports to Cloud Object Storage.
- Automated 10+ operational workflows and monitoring tasks using shell scripting, saving 30+ hours per week and reducing operational errors by 30%.
- Currently, leading a team of 6 engineers, conducting knowledge transfers, code reviews, mentoring two interns, and driving delivery priorities to improve onboarding speed and team productivity.

Software Developer

Jan 2023 – Apr

2025

- Built Excel-driven batch automation for Sales role creation, reducing processing time from minutes per role to minutes for hundreds of roles (99% time reduction). Project won 2nd Place in 1Q25 Client Advocacy Awards.
- Developed user-centric web applications with clear, context-aware exception handling, reducing user errors by 30% and improving overall UX.
- Created a natural-language chatbot integrating IBM Db2 with LLM summarization to deliver real-time system health insights, improving incident response times by 40%.

- Built an AI-assisted incident resolution platform leveraging historical data to suggest probable root causes and fixes, improving resolution efficiency by 25%.
- Engineered a prompt-driven GenAI system to auto-generate unit and integration test cases, increasing test coverage by 30% and reducing manual QA effort.
- Led an 11-member team to build a Judges' Award-winning multi-agent LLM system at IBM Developers Jumpstart, reducing manual code review effort by 60% and accelerating approval cycles.

Project Work

LLM-Powered Multimodal Quiz Generator

- Built an end-to-end multimodal quiz generation platform supporting text, documents, images, audio, video, and URLs via a unified processing pipeline.
- Implemented adaptive MCQ generation and validation using LLMs, OCR, audio/video transcription, and modular pipeline architecture.
- Mistral LLM, Prompt Engineering, OCR, Whisper (Audio-to-Text), MoviePy (Video-to-Text), PyMuPDF, python-docx, BeautifulSoup, Python, Modular Pipeline Architecture, Streamlit.

AI Portfolio Assistant (RAG-Based)

- Integrated a RAG-powered AI assistant into a personal portfolio, enabling natural-language exploration of content with strict grounding rules to prevent hallucinations.
- Deployed locally and on cloud using Flask, vector databases, and LLM inference pipelines for reliable recruiter-facing interactions.
- Llama 3.1 (8B), Sentence Transformers (all-MiniLM-L6-v2), RAG Architecture, Python, Flask, FlaskCORS, ChromaDB, Ollama (Local deployment), Playwright, React, JavaScript, CSS3, Render, Vercel.

Education

Ramachandra College of Engineering, Eluru | JNTU-Kakinada

Aug 2019 – May 2023

B.Tech in Computer Science and Engineering

CGPA: 8.14/10

Relevant Coursework: Object Oriented Programming, Databases, Discrete Maths, Data Structures and Algorithms, Operating Systems, Computer Networks, Machine Learning, Deep Learning, Artificial Intelligence, Data Mining, NLP, Advanced Data Structures and Algorithms, Information Retrieval, Image Processing.

Awards and Recognition

- **1st Place & Judges' Award** - IBM Developers Jumpstart (Multi-Agent LLM System).
- **2nd Place** - 1Q25 Client Advocacy Awards (Sales Plan Automation).

Early Projects and Certifications

- Named Entity Recognition using Bi-LSTM.
- Machine Learning-Based Housing Price Prediction.
- Certified Machine Learning Intern at Pantechelearning.